


## ECE1755 Assignment-1 Report

In this lab, I have implemented six optimization strategies to speed up the computation. The baseline serial version runs for approximately **17 secs**.

### 1. Thread-level Parallelization

This approach speeds up the execution most significantly. The speedup is **4.5 ×**. The workload is segmented into trunks based on the number of threads, as shown in Fig. 1.



In such a way, there is no synchronization required on the thread level. The accesses to tx data (point\_x, point\_y, point\_z), rx\_data and image are all independent.

### 2. Merge Tx-Rx Computation Loops and Re-use Duplicated Variables

This approach is implemented right after threads parallelization, and the speedup is on the **millisecond** scale ( $\sim 0.5\text{ ms}$ ), pushing the overall speedup to **5 ×**. Hence, it makes the memory access more compact.

### 3. Intel Intrinsic Vector Operation

Inspired by Piazza post and lecture note, single instruction and multiple data execution accelerated the program by 1~1.5 secs. The overall speedup is around **6~6.5 ×**.

### 4. GCC Ofast Compilation Flag

For this computation intense program, `-ffast-math` flag included in GCC is shown not causing any floating-number calculation. we assume that that *all math is finite*, no need to check NaN (or zero). It is also safe to disable `errno`, which saves one writing to a thread-local variable. Since we index `rx_data` with `int` and divide a constant, the enabled *reciprocal approximations* for division and reciprocal square root also proves to work.

### 5. restrict Keyword

Opposite to what `volatile` indicates the compiler to fetch from memory every time using a variable, the `restrict` C-only keyword tells the compiler that for the lifetime of this pointer, no other pointers will be used to access this specific object, i.e. an array of `float`.

### 6. Single write at the end of image shading loop

Instead of writing directly to the `image` for the approximately 1000 times at each loop iteration, we keep the temporary sum in a float and perform a single write at the end of this loop.

Approaches 4-6 jointly bring the speedup to be around **7 ×**, around **2.5~2.6 secs**.